

## ✓ Lab 6 : RNN, LSTM & Transformers

1. Write a Python program to implement a text classification model using RNN and LSTM on the Twitter Sentiment140 dataset.

- Use pre-trained Word2Vec embeddings.
- Train both RNN and LSTM models.
- Compare their accuracy and training time.

```
#Importing required libraries
import pandas as pd
import re

dataset = pd.read_csv("twitter_sentiment.csv")
dataset.head()
```



	text	label
0	Feeling great today! Life is good.	1
1	Ugh, this weather sucks!	0
2	Can't wait for the weekend!	1
3	Traffic jam ruined my morning.	0
4	Loving the new phone!	1

## ✓ Data Preprocessing

Some General Preprocessing and based on the assumptions for Twitter reviews. Rest have to check few records to find the pattern.

- Remove Null Values
- Remove punctuations
- Lower Case
- Remove digits and words containing digits
- Remove urls, mentions (@username) and topics (#viral)
- Remove stop words
- Lemmatization

```
# Display random sample of tweets
print(dataset.sample(10))
```



	text	label
44	Sunshine and smiles!	1
48	Lovely dinner with family ❤️	1
27	Missed my deadline 😞	0
45	Got ignored all day.	0
24	Just finished a good book.	1
37	Too much work, no time to rest.	0
46	My favorite band released a new song!	1
7	Horrible service at the cafe today.	0
31	Spilled coffee all over myself.	0
30	Woke up early and feeling fresh!	1

```
# Display example tweets containing URLs, mentions, hashtags, emojis, etc.
sample_tweets = dataset[
    dataset['text'].str.contains(r"http|@|#", regex=True, na=False)
]

print(sample_tweets)
```



```
Empty DataFrame
Columns: [text, label]
Index: []
```

- There are no urls, mentions and hashtags
- But we have emojis

```
import nltk
import string
from nltk.corpus import stopwords
from nltk.tokenize import word_tokenize
```

```
from nltk.stem import WordNetLemmatizer
import emoji
```

```
!pip install emoji
```

```
Collecting emoji
  Downloading emoji-2.14.1-py3-none-any.whl.metadata (5.7 kB)
  Downloading emoji-2.14.1-py3-none-any.whl (590 kB)
    ----- 0.0/590.6 kB ? eta -:--:--
    ----- 262.1/590.6 kB ? eta -:--:--
    ----- 590.6/590.6 kB 1.2 MB/s eta 0:00:00
Installing collected packages: emoji
Successfully installed emoji-2.14.1
```

```
# Download necessary NLTK resources (only once)
nltk.download('stopwords')
nltk.download('punkt')
nltk.download('wordnet')
nltk.download('punkt_tab')
```

```
[nltk_data] Downloading package stopwords to
[nltk_data] C:\Users\alens\AppData\Roaming\nltk_data...
[nltk_data] Package stopwords is already up-to-date!
[nltk_data] Downloading package punkt to
[nltk_data] C:\Users\alens\AppData\Roaming\nltk_data...
[nltk_data] Package punkt is already up-to-date!
[nltk_data] Downloading package wordnet to
[nltk_data] C:\Users\alens\AppData\Roaming\nltk_data...
[nltk_data] Package wordnet is already up-to-date!
[nltk_data] Downloading package punkt_tab to
[nltk_data] C:\Users\alens\AppData\Roaming\nltk_data...
[nltk_data] Unzipping tokenizers\punkt_tab.zip.
True
```

```
stop_words = set(stopwords.words('english'))
lemmatizer = WordNetLemmatizer()

def remove_emojis(text):
    return emoji.replace_emoji(text, replace='')

def preprocess_text(text):
    text = text.lower() #lowercase
    text = remove_emojis(text) # remove emoji
    text = re.sub(r'[\w\s]', '', text) # remove punctuations
    text = re.sub(r'\b\w*\d\w*\b', '', text) #remove digits and words with digits
    words = word_tokenize(text)

    words = [lemmatizer.lemmatize(word) for word in words if word not in stop_words]

    return ' '.join(words)

dataset['clean_text'] = dataset['text'].apply(preprocess_text)

(dataset.head(10))
```

	text	label	clean_text
0	Feeling great today! Life is good.	1	feeling great today life good
1	Ugh, this weather sucks!	0	ugh weather suck
2	Can't wait for the weekend!	1	cant wait weekend
3	Traffic jam ruined my morning.	0	traffic jam ruined morning
4	Loving the new phone!	1	loving new phone
5	Totally exhausted and frustrated.	0	totally exhausted frustrated
6	My dog just learned a new trick! So proud.	1	dog learned new trick proud
7	Horrible service at the cafe today.	0	horrible service cafe today
8	Had an amazing run. Endorphins are real!	1	amazing run endorphin real
9	Missed my flight. Worst day ever.	0	missed flight worst day ever

## Use Pretrained Word2Vec Embedding

```
import gensim.downloader as api
```

```
print(list(api.info()['models'].keys()))
```

```
↳ ['fasttext-wiki-news-subwords-300', 'conceptnet-numberbatch-17-06-300', 'word2vec-ruscorpora-300', 'word2vec-google-news-300', 'glove-twitter-50']
```

```
word2vec_model = api.load("glove-twitter-50")
```

```
↳ [=====] 100.0% 199.5/199.5MB downloaded
```

```
max_len = dataset['clean_text'].str.len().max()
print(max_len)
```

```
↳ 33
```

```
import numpy as np
from tensorflow.keras.preprocessing.text import Tokenizer
from tensorflow.keras.preprocessing.sequence import pad_sequences
```

```
# Tokenize words
tokenizer = Tokenizer()
tokenizer.fit_on_texts(dataset['clean_text'])
```

```
# Convert text to sequences
sequences = tokenizer.texts_to_sequences(dataset['clean_text'])
```

```
# Set a fixed length for sequences
MAX_LEN = 30
padded_sequences = pad_sequences(sequences, maxlen=MAX_LEN, padding='post')
```

```
# Get vocabulary size
vocab_size = len(tokenizer.word_index) + 1
```

```
# Create Embedding Matrix
```

```
EMBEDDING_DIM = 50 #dim of glove-twitter-50
```

```
# Initialize an embedding matrix
embedding_matrix = np.zeros((vocab_size, EMBEDDING_DIM))
```

```
for word, i in tokenizer.word_index.items():
    if word in word2vec_model:
        embedding_matrix[i] = word2vec_model[word]
```

## ✓ RNN Model

```
import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Embedding, SimpleRNN, Dense, Dropout
```

```
model = Sequential([
    Embedding(input_dim=vocab_size, output_dim=EMBEDDING_DIM,
              weights=[embedding_matrix], trainable=False),
    SimpleRNN(128, activation='tanh', return_sequences=True),
    SimpleRNN(64, activation='tanh'),
    Dropout(0.5),
    Dense(1, activation='sigmoid')
])
```

```
# Compile model
model.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])
```

```
# Model summary
model.summary()
```

Model: "sequential\_1"

Layer (type)	Output Shape	Param #
embedding_1 (Embedding)	?	6,200
simple_rnn_2 (SimpleRNN)	?	0 (unbuilt)
simple_rnn_3 (SimpleRNN)	?	0 (unbuilt)
dropout_1 (Dropout)	?	0
dense_1 (Dense)	?	0 (unbuilt)

Total params: 6,200 (24.22 KB)

```
from sklearn.model_selection import train_test_split
import time

X_train, X_test, y_train, y_test = train_test_split(padded_sequences, dataset['label'], test_size=0.2, random_state=42)

# Train model
start_time = time.time()
history = model.fit(X_train, y_train, epochs=10, batch_size=32, validation_data=(X_test, y_test))
end_time = time.time()

training_time = end_time - start_time
print(f"Training Time: {training_time:.2f} seconds")
```

```
Epoch 1/10
2/2 ————— 6s 859ms/step - accuracy: 0.5750 - loss: 0.7311 - val_accuracy: 0.7000 - val_loss: 0.5910
Epoch 2/10
2/2 ————— 0s 139ms/step - accuracy: 0.6958 - loss: 0.6091 - val_accuracy: 0.9000 - val_loss: 0.5005
Epoch 3/10
2/2 ————— 0s 218ms/step - accuracy: 0.7333 - loss: 0.5270 - val_accuracy: 0.8000 - val_loss: 0.4819
Epoch 4/10
2/2 ————— 0s 144ms/step - accuracy: 0.8375 - loss: 0.4164 - val_accuracy: 0.7000 - val_loss: 0.5618
Epoch 5/10
2/2 ————— 0s 140ms/step - accuracy: 0.9125 - loss: 0.3425 - val_accuracy: 0.7000 - val_loss: 0.6061
Epoch 6/10
2/2 ————— 0s 141ms/step - accuracy: 0.9563 - loss: 0.2310 - val_accuracy: 0.7000 - val_loss: 0.5063
Epoch 7/10
2/2 ————— 0s 130ms/step - accuracy: 0.9458 - loss: 0.2354 - val_accuracy: 0.8000 - val_loss: 0.4399
Epoch 8/10
2/2 ————— 0s 140ms/step - accuracy: 0.9729 - loss: 0.1778 - val_accuracy: 0.9000 - val_loss: 0.3934
Epoch 9/10
2/2 ————— 0s 190ms/step - accuracy: 0.9458 - loss: 0.1407 - val_accuracy: 0.9000 - val_loss: 0.3525
Epoch 10/10
2/2 ————— 0s 145ms/step - accuracy: 0.9729 - loss: 0.0890 - val_accuracy: 0.8000 - val_loss: 0.3407
Training Time: 8.13 seconds
```

```
final_train_acc = history.history['accuracy'][-1]
final_val_acc = history.history['val_accuracy'][-1]

print(f"Final Training Accuracy: {final_train_acc:.4f}")
print(f"Final Validation Accuracy: {final_val_acc:.4f}")
```

```
Final Training Accuracy: 0.9750
Final Validation Accuracy: 0.8000
```

## ✓ LSTM Model

```
from tensorflow.keras.layers import LSTM

# Define the LSTM model
lstm_model = Sequential([
    Embedding(input_dim=vocab_size, output_dim=50,
              weights=[embedding_matrix], trainable=False),
    LSTM(64, return_sequences=False),
    Dropout(0.5),
    Dense(1, activation='sigmoid')
])

# Compile the model
lstm_model.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])

# Summary
lstm_model.summary()
```

Model: "sequential\_2"

Layer (type)	Output Shape	Param #
embedding_2 (Embedding)	?	6,200
lstm (LSTM)	?	0 (unbuilt)
dropout_2 (Dropout)	?	0
dense_2 (Dense)	?	0 (unbuilt)

Total params: 6,200 (24.22 KB)

Trainable params: 0 (0.00 KB)

```
start_time = time.time()
lstm_history = lstm_model.fit(X_train, y_train, epochs=10, batch_size=32, validation_data=(X_test, y_test))
end_time = time.time()
```

```
lstm_training_time = end_time - start_time
print(f"LSTM Training Time: {lstm_training_time:.2f} seconds")
```

```
final_lstm_train_acc = lstm_history.history['accuracy'][-1]
final_lstm_val_acc = lstm_history.history['val_accuracy'][-1]
```

```
print(f"LSTM Final Training Accuracy: {final_lstm_train_acc:.4f}")
print(f"LSTM Final Validation Accuracy: {final_lstm_val_acc:.4f}")
```

```
Epoch 1/10
2/2 ————— 5s 730ms/step - accuracy: 0.5437 - loss: 0.6932 - val_accuracy: 0.4000 - val_loss: 0.6934
Epoch 2/10
2/2 ————— 0s 140ms/step - accuracy: 0.5333 - loss: 0.6925 - val_accuracy: 0.4000 - val_loss: 0.6940
Epoch 3/10
2/2 ————— 0s 136ms/step - accuracy: 0.5479 - loss: 0.6910 - val_accuracy: 0.4000 - val_loss: 0.6943
Epoch 4/10
2/2 ————— 0s 150ms/step - accuracy: 0.5167 - loss: 0.6903 - val_accuracy: 0.4000 - val_loss: 0.6938
Epoch 5/10
2/2 ————— 0s 150ms/step - accuracy: 0.5271 - loss: 0.6889 - val_accuracy: 0.4000 - val_loss: 0.6927
Epoch 6/10
2/2 ————— 0s 140ms/step - accuracy: 0.5062 - loss: 0.6859 - val_accuracy: 0.4000 - val_loss: 0.6904
Epoch 7/10
2/2 ————— 0s 139ms/step - accuracy: 0.5708 - loss: 0.6769 - val_accuracy: 0.4000 - val_loss: 0.6864
Epoch 8/10
2/2 ————— 0s 143ms/step - accuracy: 0.5646 - loss: 0.6702 - val_accuracy: 0.4000 - val_loss: 0.6777
Epoch 9/10
2/2 ————— 0s 144ms/step - accuracy: 0.5604 - loss: 0.6538 - val_accuracy: 0.4000 - val_loss: 0.6597
Epoch 10/10
2/2 ————— 0s 137ms/step - accuracy: 0.5979 - loss: 0.6042 - val_accuracy: 0.5000 - val_loss: 0.6351
LSTM Training Time: 6.70 seconds
LSTM Final Training Accuracy: 0.6000
LSTM Final Validation Accuracy: 0.5000
```

# Compare results

```
print(f"\nRNN Training Time: {training_time:.2f} sec vs LSTM Training Time: {lstm_training_time:.2f} sec")
print(f"RNN Validation Accuracy: {final_val_acc:.4f} vs LSTM Validation Accuracy: {final_lstm_val_acc:.4f}")
```

```
RNN Training Time: 8.13 sec vs LSTM Training Time: 6.70 sec
RNN Validation Accuracy: 0.8000 vs LSTM Validation Accuracy: 0.5000
```

2. Write a Python program using a pre-trained Transformer model (like BART or T5) from Hugging Face to perform text summarization.

- Accept a long article as input.
- Generate a concise summary.
- Evaluate using the different evaluation matrix scores.

```
!pip install torch transformers
```



```

Attempting uninstall: nvidia-nvjitlink-cu12
Found existing installation: nvidia-nvjitlink-cu12 12.5.82
Uninstalling nvidia-nvjitlink-cu12-12.5.82:
Successfully uninstalled nvidia-nvjitlink-cu12-12.5.82
Attempting uninstall: nvidia-curand-cu12
Found existing installation: nvidia-curand-cu12 10.3.6.82
Uninstalling nvidia-curand-cu12-10.3.6.82:
Successfully uninstalled nvidia-curand-cu12-10.3.6.82
Attempting uninstall: nvidia-cufft-cu12
Found existing installation: nvidia-cufft-cu12 11.2.3.61
Uninstalling nvidia-cufft-cu12-11.2.3.61:
Successfully uninstalled nvidia-cufft-cu12-11.2.3.61
Attempting uninstall: nvidia-cuda-runtime-cu12
Found existing installation: nvidia-cuda-runtime-cu12 12.5.82
Uninstalling nvidia-cuda-runtime-cu12-12.5.82:
Successfully uninstalled nvidia-cuda-runtime-cu12-12.5.82
Attempting uninstall: nvidia-cuda-nvrtc-cu12
Found existing installation: nvidia-cuda-nvrtc-cu12 12.5.82
Uninstalling nvidia-cuda-nvrtc-cu12-12.5.82:
Successfully uninstalled nvidia-cuda-nvrtc-cu12-12.5.82
Attempting uninstall: nvidia-cuda-cupti-cu12
Found existing installation: nvidia-cuda-cupti-cu12 12.5.82
Uninstalling nvidia-cuda-cupti-cu12-12.5.82:
Successfully uninstalled nvidia-cuda-cupti-cu12-12.5.82
Attempting uninstall: nvidia-cublas-cu12
Found existing installation: nvidia-cublas-cu12 12.5.3.2
Uninstalling nvidia-cublas-cu12-12.5.3.2:
Successfully uninstalled nvidia-cublas-cu12-12.5.3.2
Attempting uninstall: nvidia-cusparse-cu12
Found existing installation: nvidia-cusparse-cu12 12.5.1.3
Uninstalling nvidia-cusparse-cu12-12.5.1.3:
Successfully uninstalled nvidia-cusparse-cu12-12.5.1.3
Attempting uninstall: nvidia-cudnn-cu12
Found existing installation: nvidia-cudnn-cu12 9.3.0.75
Uninstalling nvidia-cudnn-cu12-9.3.0.75:
Successfully uninstalled nvidia-cudnn-cu12-9.3.0.75
Attempting uninstall: nvidia-cusolver-cu12
Found existing installation: nvidia-cusolver-cu12 11.6.3.83
Uninstalling nvidia-cusolver-cu12-11.6.3.83:
Successfully uninstalled nvidia-cusolver-cu12-11.6.3.83
Successfully installed nvidia-cublas-cu12-12.4.5.8 nvidia-cuda-cupti-cu12-12.4.127 nvidia-cuda-nvrtc-cu12-12.4.127 nvidia-cuda-ru

```

```


import torch
from transformers import AutoTokenizer, AutoModelWithLMHead

```

```

# T5: Text To Text Trasfer Transformer
tokenizer=AutoTokenizer.from_pretrained('T5-base')
model=AutoModelWithLMHead.from_pretrained('T5-base', return_dict=True)

```


 /usr/local/lib/python3.11/dist-packages/huggingface\_hub/utils/\_auth.py:94: UserWarning:  
The secret `HF\_TOKEN` does not exist in your Colab secrets.  
To authenticate with the Hugging Face Hub, create a token in your settings tab (<https://huggingface.co/settings/tokens>), set it as :  
You will be able to reuse this secret in all of your notebooks.  
Please note that authentication is recommended but still optional to access public models or datasets.

```

warnings.warn(
config.json: 100% ██████████ 1.21k/4.21k [00:00<00:00, 37.9kB/s]
spiece.model: 100% ██████████ 792k/792k [00:00<00:00, 1.26MB/s]
tokenizer.json: 100% ██████████ 1.39M/1.39M [00:00<00:00, 1.75MB/s]
/usr/local/lib/python3.11/dist-packages/transformers/models/auto/modeling_auto.py:1881: FutureWarning: The class `AutoModelWithLMHead`
warnings.warn(
model.safetensors: 100% ██████████ 892M/892M [00:05<00:00, 167MB/s]
generation_config.json: 100% ██████████ 147/147 [00:00<00:00, 13.4kB/s]

```

```
para = input("Enter a Para to be summarized : ")
```

 Enter a Para to be summarized : Ad sales boost Time Warner profit Quarterly profits at US media giant TimeWarner jumped 76% to \$1.1

```

inputs=tokenizer.encode("sumarize: " +para,return_tensors='pt', max_length=512, truncation=True)

output = model.generate(inputs, min_length=80, max_length=200)

summary=tokenizer.decode(output[0])

```

```

import textwrap

# Wrap the summary text to a fixed width (e.g., 80 characters per line)
wrapped_summary = textwrap.fill(summary, width=80)

```

```
print(wrapped_summary)
```



```
<pad> time warner's profit jumped 76% to $1.13bn (£600m) for the three months to  
December . firm benefited from sales of high-speed internet connections and  
higher advert sales . time warner's film division saw profits slump 27% to $284m  
. it is to restate its accounts as part of efforts to resolve an inquiry into  
AOL .</s>
```

```
!pip install rouge-score
```



```
Collecting rouge-score  
  Downloading rouge_score-0.1.2.tar.gz (17 kB)  
  Preparing metadata (setup.py) ... done  
Requirement already satisfied: absl-py in /usr/local/lib/python3.11/dist-packages (from rouge-score) (1.4.0)  
Requirement already satisfied: nltk in /usr/local/lib/python3.11/dist-packages (from rouge-score) (3.9.1)  
Requirement already satisfied: numpy in /usr/local/lib/python3.11/dist-packages (from rouge-score) (2.0.2)  
Requirement already satisfied: six>=1.14.0 in /usr/local/lib/python3.11/dist-packages (from rouge-score) (1.17.0)  
Requirement already satisfied: click in /usr/local/lib/python3.11/dist-packages (from nltk->rouge-score) (8.1.8)  
Requirement already satisfied: joblib in /usr/local/lib/python3.11/dist-packages (from nltk->rouge-score) (1.4.2)  
Requirement already satisfied: regex>=2021.8.3 in /usr/local/lib/python3.11/dist-packages (from nltk->rouge-score) (2024.11.6)  
Requirement already satisfied: tqdm in /usr/local/lib/python3.11/dist-packages (from nltk->rouge-score) (4.67.1)  
Building wheels for collected packages: rouge-score  
  Building wheel for rouge-score (setup.py) ... done  
  Created wheel for rouge-score: filename=rouge_score-0.1.2-py3-none-any.whl size=24935 sha256=776758668c7c69580f5f83b9f96de2cd60cb6  
  Stored in directory: /root/.cache/pip/wheels/1e/19/43/8a442dc83660ca25e163e1bd1f89919284ab0d0c1475475148  
Successfully built rouge-score  
Installing collected packages: rouge-score  
Successfully installed rouge-score-0.1.2
```

```
from rouge_score import rouge_scorer  
  
def evaluate_summary(reference, generated):  
    scorer = rouge_scorer.RougeScorer(['rouge1', 'rouge2', 'rougeL'], use_stemmer=True)  
    scores = scorer.score(reference, generated)  
    return scores  
  
reference_summary = input("Enter the Reference Summary: ")  
  
rouge_scores = evaluate_summary(reference_summary, summary)  
  
print("\nEvaluation Metrics:")  
for key, value in rouge_scores.items():  
    print(f"{key.upper()}: Precision={value.precision:.4f}, Recall={value.recall:.4f}, F1-score={value.fmeasure:.4f}")
```



```
Enter the Reference Summary: TimeWarner said fourth quarter sales rose 2% to $11.1bn from $10.9bn.For the full-year, TimeWarner post  
  
Evaluation Metrics:  
ROUGE1: Precision=0.5690, Recall=0.2157, F1-score=0.3128  
ROUGE2: Precision=0.2456, Recall=0.0921, F1-score=0.1340  
ROUGEL: Precision=0.3793, Recall=0.1438, F1-score=0.2085
```